

# Figure captions

Captions copied verbatim from `manuscript/drafts/v7_draft.tex`. Image files are in `manuscript/final/figures/`, converted from the 600 dpi source PNGs in `manuscript/figures/` (generated by `manuscript/figures/make_figures.py`) to TIFF at their native 600 dpi — a format conversion, not a resolution upscale. 600 dpi exceeds JPD's 300 dpi color minimum for all four (these are color/mixed raster figures, not pure line art or grayscale, so the 300 dpi threshold applies).

## Figure 1 (`Figure1.tif`)

Sample construction and role structure. Of 100 submissions, 13 were abandoned before the role fork and 14 of the remaining 87 screened out by answering "no" to both role items, leaving an analytic sample of  $N=73$ . Roles are not mutually exclusive: the 7 dual-role respondents enter both branches, giving 72 participant-side and 8 facilitator-side reports, and 3 of those 7 were stopped by the participant continuation gate and 2 further participants answered "prefer not to answer" to the facilitator role item, satisfying neither arm of the participant section's relevance condition, so 67 respondents were actually routed to the participant items. The branches are unlinked at the ceremony level — no identifier connects a participant to the facilitator who ran their ceremony, and the ceremonies a dual-role respondent led are different events from those they attended — so no participant–facilitator comparison is possible.

## Figure 2 (`Figure2.tif`)

Substances used across all ceremonies attended, participant-reported (select-multiple; base  $n=67$  participants answering). Bars are labeled with the count and the percentage of that base; because respondents could select more than one substance, percentages need not sum to 100%. The "other" bar is drawn in grey because its free-text follow-up was completed by only 3 respondents. Distinct from the facilitator-led substance variable and from the most-recent-ceremony substance variable.

## Figure 3 (`Figure3.tif`)

Distribution of reported satisfaction with the most recent ceremony (0 = very unpleasant to 6 = very pleasant). Of the 72 respondents identifying as participants, 63 gave a numeric score and form the base; the remaining 9 comprise 1 "prefer not to answer", 5 structural skips, and 3 item nonresponses.

Median 6, with 51 of 63 (81%) answering 5 or 6 — the ceiling that limits how much variance any predictor can explain.

## Figure 4 (`Figure4.tif`)

Reported satisfaction by self-directed preparation status. Any preparation,  $n=58$  answering the outcome (median 6, IQR 5–6); no preparation,  $n=5$  (median 4, IQR 3–5). Marker area is proportional to the number of respondents giving each score, since scores tie heavily at the ceiling. The no-preparation group is drawn as its five raw points and a median tick rather than a box, which at that  $n$  would imply distributional precision the data cannot support. **Exploratory:** this comparison is suggestive, not confirmatory.

## Note on the graphical abstract

`manuscript/figures/graphical_abstract.png` exists in the project but has **no corresponding figure environment, caption, or `\includegraphics` call in `v7_draft.tex`** — it is not part of the approved manuscript and JPD's own instructions page does not list a graphical abstract as a submission element (confirmed in the original `JOURNAL_REQUIREMENTS.md`; not part of this pass's re-verification, but nothing in the 2026-09-08 re-check found a graphical-abstract requirement either). It is therefore **not** included in `manuscript/final/figures/` or the submission package.